

Ансамбли моделей машинного обучения

In [70]:

```
import numpy as np
import pandas as pd
from typing import Dict, Tuple
from scipy import stats
from scipy.optimize import fmin_tnc
from IPython.display import Image
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_squared_log_error, median_absolute_error, r2_score, root_mean_squared_error
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import SGDRegressor
from sklearn.linear_model import SGDClassifier
from sklearn import linear_model
from sklearn.ensemble import BaggingRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import ExtraTreesRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
import xgboost as xgb
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
```

In [6]:

```
mpg = pd.read_csv('C:\\MGTU\\6 semestr\\TMO\\auto-mpg.csv')
```

In [7]:

```
mpg.head()
```

Out[7]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino

In [8]:

```
mpg.dtypes
```

Out[8]:

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    object
weight        int64
acceleration  float64
model year    int64
```

```
model year      int64
origin          int64
car name        object
dtype: object
```

In [9]:

```
mpg = mpg[mpg['horsepower'] != '?']
mpg['horsepower'] = mpg['horsepower'].astype(float)
```

In [10]:

```
mpg.dtypes
```

Out[10]:

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    float64
weight        int64
acceleration  float64
model year    int64
origin        int64
car name      object
dtype: object
```

In [11]:

```
mpg = mpg.drop(columns=['car name'])
```

In [12]:

```
X = mpg.drop(columns=['mpg']) # Признаки
y = mpg['mpg'] # Целевая переменная

# Разделение данных на обучающую и тестовую выборки
mpg_X_train, mpg_X_test, mpg_y_train, mpg_y_test = train_test_split(X, y, test_size=0.2,
random_state=1)
```

Bagging

In [53]:

```
br_mpg = BaggingRegressor(n_estimators=1000, oob_score = False, random_state=10)
br_mpg.fit(mpg_X_train,mpg_y_train )
br_mpg_predict = br_mpg.predict(mpg_X_test)
```

In [54]:

```
root_mean_squared_error(mpg_y_test, br_mpg_predict)
```

Out[54]:

```
3.1432575417344863
```

Random Forest

In [55]:

```
rfr_mpg = RandomForestRegressor(n_estimators=1000, oob_score = True, random_state=10)
rfr_mpg.fit(mpg_X_train,mpg_y_train )
rfr_mpg_predict = rfr_mpg.predict(mpg_X_test)
```

In [56]:

```
root_mean_squared_error(mpg_y_test, rfr_mpg_predict)
```

Out[56]:

3.151838599486819

In [89]:

```
rfr_mpg1 = RandomForestRegressor(n_estimators=1000, oob_score = True, max_features = 2 ,
min_samples_leaf = 5, random_state=10)
rfr_mpg1.fit(mpg_X_train,mpg_y_train )
rfr_mpg_predict1 = rfr_mpg1.predict(mpg_X_test)
```

In [90]:

```
root_mean_squared_error(mpg_y_test, rfr_mpg_predict1)
```

Out[90]:

3.55224352114633

Extremely Randomized Trees

In [59]:

```
etr_mpg = ExtraTreesRegressor(n_estimators=1000, random_state=10)
etr_mpg.fit(mpg_X_train,mpg_y_train )
etr_mpg_predict = etr_mpg.predict(mpg_X_test)
```

In [60]:

```
root_mean_squared_error(mpg_y_test, etr_mpg_predict)
```

Out[60]:

3.312486072538161

AdaBoost

In [64]:

```
abr_mpg = AdaBoostRegressor(n_estimators=1000, random_state=10)
abr_mpg.fit(mpg_X_train,mpg_y_train )
abr_mpg_predict = abr_mpg.predict(mpg_X_test)
```

In [65]:

```
root_mean_squared_error(mpg_y_test, abr_mpg_predict)
```

Out[65]:

3.620154134360523

GradientBoosting

In [67]:

```
gbr_mpg = GradientBoostingRegressor(n_estimators=1000, random_state=10)
gbr_mpg.fit(mpg_X_train,mpg_y_train )
gbr_mpg_predict = gbr_mpg.predict(mpg_X_test)
```

In [68]:

```
root_mean_squared_error(mpg_y_test, gbr_mpg_predict)
```

Out[68]:

3.1436494559859263

XGBoost

In [77]:

```
dtrain = xgb.DMatrix(mpg_X_train, label=mpg_y_train)
dtest = xgb.DMatrix(mpg_X_test, label=mpg_y_test)

params = {
    'objective': 'reg:squarederror', # функция потерь для задачи регрессии
    'eval_metric': 'rmse'
}

num_rounds = 1000
model = xgb.train(params, dtrain, num_rounds)

y_pred = model.predict(dtest)
```

In [78]:

```
root_mean_squared_error(mpg_y_test, y_pred)
```

Out[78]:

3.2797096825132277